

# Continuous\_Control

August 5, 2020

## 1 Continuous Control

---

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

### 1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible.  
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.0.
```

The environments corresponding to both versions of the environment are already saved in the Workspace and can be accessed at the file paths provided below.

Please select one of the two options below for loading the environment.

```
In [2]: from unityagents import UnityEnvironment  
import numpy as np  
  
# select this option to load version 1 (with a single agent) of the environment  
env = UnityEnvironment(file_name='/data/Reacher_One_Linux_NoVis/Reacher_One_Linux_NoVis.  
  
# select this option to load version 2 (with 20 agents) of the environment  
# env = UnityEnvironment(file_name='/data/Reacher_Linux_NoVis/Reacher.x86_64')
```

```
INFO:unityagents:  
'Academy' started successfully!  
Unity Academy name: Academy  
Number of Brains: 1  
Number of External Brains : 1  
Lesson number : 0  
Reset Parameters :  
    goal_speed -> 1.0
```

```

        goal_size -> 5.0
Unity brain name: ReacherBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 33
    Number of stacked Vector Observation: 1
    Vector Action space type: continuous
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

## 1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents
        num_agents = len(env_info.agents)
        print('Number of agents:', num_agents)

        # size of each action
        action_size = brain.vector_action_space_size
        print('Size of each action:', action_size)

        # examine the state space
        states = env_info.vector_observations
        state_size = states.shape[1]
        print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
        print('The state for the first agent looks like:', states[0])

```

```

Number of agents: 1
Size of each action: 4
There are 1 agents. Each observes a state with length: 33
The state for the first agent looks like: [ 0.00000000e+00 -4.00000000e+00  0.00000000e+00
-0.00000000e+00 -0.00000000e+00 -4.37113883e-08  0.00000000e+00
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00  0.00000000e+00 -1.00000000e+01  0.00000000e+00
 1.00000000e+00 -0.00000000e+00 -0.00000000e+00 -4.37113883e-08
 0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00]

```

```

0.00000000e+00  0.00000000e+00  5.75471878e+00 -1.00000000e+00
5.55726671e+00  0.00000000e+00  1.00000000e+00  0.00000000e+00
-1.68164849e-01]

```

### 1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agents while they are training**, and you should set `train_mode=True` to restart the environment.

```

In [5]: env_info = env.reset(train_mode=True)[brain_name]           # reset the environment
        states = env_info.vector_observations                       # get the current state (for each
        scores = np.zeros(num_agents)                             # initialize the score (for each
        while True:
            actions = np.random.randn(num_agents, action_size)    # select an action (for each agent)
            actions = np.clip(actions, -1, 1)                      # all actions between -1 and 1
            env_info = env.step(actions)[brain_name]               # send all actions to the environment
            next_states = env_info.vector_observations              # get next state (for each agent)
            rewards = env_info.rewards                             # get reward (for each agent)
            dones = env_info.local_done                            # see if episode finished
            scores += env_info.rewards                             # update the score (for each agent)
            states = next_states                                    # roll over states to next time step
            if np.any(dones):                                       # exit loop if episode finished
                break
        print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores)))

```

```
Total score (averaged over agents) this episode: 0.13999999687075615
```

When finished, you can close the environment.

```
In [ ]: env.close()
```

### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agents while they are training. However, *after training the agents*, you can download the saved model weights to watch the agents on your own machine!

```
In [11]: from ddpq_agent import Agent
         from collections import deque
         import torch
```

```
In [15]: # DDPG function
```

```
def ddpq(n_episodes=1000, max_t=1000, consec_episodes=100, print_every=1):
    """Deep Deterministic Policy Gradient (DDPG)

    Params
    =====
        n_episodes (int)      : maximum number of training episodes
        max_t (int)           : maximum number of timesteps per episode
        consec_episodes (int) : number of consecutive episodes used to calculate score
        print_every (int)     : interval to display results

    """

    scores = []
    scores_deque = deque(maxlen=print_every)

    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name]
        agent.reset()
        state = env_info.vector_observations[0]           # get the current state
        score = 0

        #print('start')

        for t in range(max_t):
            action = agent.act(state, add_noise=True)    # select an action

            env_info = env.step(action)[brain_name]      # send the action to the env
            next_state = env_info.vector_observations[0] # get the next state
            reward = env_info.rewards[0]                 # get the reward
            done = env_info.local_done[0]                # see if episode has finished
            agent.step(state, action, reward, next_state, done) # take step with agent
            score += reward                               # update the score
            state = next_state                           # roll over the state to next
            if done:                                     # exit loop if episode finished
                break

        scores_deque.append(score)                       # save most recent score
        scores.append(score)                             # save most recent score
```

```

# print('end')
if i_episode % print_every == 0:
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_deque)))
    torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
    torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')

if np.mean(scores_deque) >= 30.0:
    print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_deque)))
    torch.save(agent.actor_local.state_dict(), 'checkpoint_actor.pth')
    torch.save(agent.critic_local.state_dict(), 'checkpoint_critic.pth')
    break

return scores

```

In [16]: # run the training loop

```
agent = Agent(state_size=state_size, action_size=action_size, random_seed=11)
```

In [17]: scores = ddpG()

Episode 1	Average Score: 0.00
Episode 2	Average Score: 0.32
Episode 3	Average Score: 0.00
Episode 4	Average Score: 0.33
Episode 5	Average Score: 0.80
Episode 6	Average Score: 0.00
Episode 7	Average Score: 0.00
Episode 8	Average Score: 0.76
Episode 9	Average Score: 0.26
Episode 10	Average Score: 0.00
Episode 11	Average Score: 0.93
Episode 12	Average Score: 0.40
Episode 13	Average Score: 1.21
Episode 14	Average Score: 1.31
Episode 15	Average Score: 0.49
Episode 16	Average Score: 1.06
Episode 17	Average Score: 1.21
Episode 18	Average Score: 1.29
Episode 19	Average Score: 0.10
Episode 20	Average Score: 1.96
Episode 21	Average Score: 2.41
Episode 22	Average Score: 1.62
Episode 23	Average Score: 0.94
Episode 24	Average Score: 2.00
Episode 25	Average Score: 0.92
Episode 26	Average Score: 2.86

Episode 27	Average Score: 1.55
Episode 28	Average Score: 2.95
Episode 29	Average Score: 2.49
Episode 30	Average Score: 1.90
Episode 31	Average Score: 2.24
Episode 32	Average Score: 2.72
Episode 33	Average Score: 4.87
Episode 34	Average Score: 2.72
Episode 35	Average Score: 2.78
Episode 36	Average Score: 2.50
Episode 37	Average Score: 3.77
Episode 38	Average Score: 3.69
Episode 39	Average Score: 2.96
Episode 40	Average Score: 6.06
Episode 41	Average Score: 4.31
Episode 42	Average Score: 3.17
Episode 43	Average Score: 2.15
Episode 44	Average Score: 3.86
Episode 45	Average Score: 4.60
Episode 46	Average Score: 3.36
Episode 47	Average Score: 2.50
Episode 48	Average Score: 4.37
Episode 49	Average Score: 2.35
Episode 50	Average Score: 3.19
Episode 51	Average Score: 3.89
Episode 52	Average Score: 6.74
Episode 53	Average Score: 4.66
Episode 54	Average Score: 3.46
Episode 55	Average Score: 6.17
Episode 56	Average Score: 2.35
Episode 57	Average Score: 5.41
Episode 58	Average Score: 4.47
Episode 59	Average Score: 6.83
Episode 60	Average Score: 2.56
Episode 61	Average Score: 2.56
Episode 62	Average Score: 6.95
Episode 63	Average Score: 8.61
Episode 64	Average Score: 3.39
Episode 65	Average Score: 5.62
Episode 66	Average Score: 8.78
Episode 67	Average Score: 11.38
Episode 68	Average Score: 6.07
Episode 69	Average Score: 4.34
Episode 70	Average Score: 7.05
Episode 71	Average Score: 6.25
Episode 72	Average Score: 6.44
Episode 73	Average Score: 10.91
Episode 74	Average Score: 6.67

Episode 75	Average Score: 3.98
Episode 76	Average Score: 11.59
Episode 77	Average Score: 11.16
Episode 78	Average Score: 6.70
Episode 79	Average Score: 13.85
Episode 80	Average Score: 10.89
Episode 81	Average Score: 17.97
Episode 82	Average Score: 10.97
Episode 83	Average Score: 19.06
Episode 84	Average Score: 14.52
Episode 85	Average Score: 16.42
Episode 86	Average Score: 27.67
Episode 87	Average Score: 17.28
Episode 88	Average Score: 14.54
Episode 89	Average Score: 18.58
Episode 90	Average Score: 17.96
Episode 91	Average Score: 11.15
Episode 92	Average Score: 15.24
Episode 93	Average Score: 17.25
Episode 94	Average Score: 12.67
Episode 95	Average Score: 19.33
Episode 96	Average Score: 19.20
Episode 97	Average Score: 21.96
Episode 98	Average Score: 20.34
Episode 99	Average Score: 17.58
Episode 100	Average Score: 11.89
Episode 101	Average Score: 24.89
Episode 102	Average Score: 23.03
Episode 103	Average Score: 23.67
Episode 104	Average Score: 14.87
Episode 105	Average Score: 15.36
Episode 106	Average Score: 24.36
Episode 107	Average Score: 23.86
Episode 108	Average Score: 26.28
Episode 109	Average Score: 25.69
Episode 110	Average Score: 31.97


Environment solved in 110 episodes!

Average Score: 31.97

```
In [18]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='DDPG')
```

```
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left');
plt.show()
```



```
In [19]: env.close()
```

```
In [ ]:
```